




Deep Computer Vision

Alexander Amini

MIT Introduction to Deep Learning

January 6, 2026




MIT Introduction to Deep Learning
introtodeeplearning.com @MITDeepLearning





**“To know what is
where by looking.”**

To discover from images what is present in the world, where things are, what actions are taking place, to predict and anticipate events in the world



The rise and impact of computer vision


Robotics



Accessibility




Biology & Medicine





Autonomous driving




Mobile computing




Impact: Facial Detection & Recognition




Impact: Self-Driving Cars



Impact: Medicine, Biology, Healthcare



Breast cancer




COVID-19



Skin cancer




Impact: Accessibility




What Computers “See”

Images are Numbers




101
99
6
5
?

Images are Numbers



167	153	174	168	160	152	129	151	172	161	155	166
165	162	169	74	76	62	33	17	110	210	180	154
160	160	50	14	34	6	10	33	45	105	159	187
206	109	5	124	131	111	120	204	165	15	56	180
194	63	137	251	237	239	239	228	227	37	71	201
172	165	297	233	233	214	220	239	228	94	74	206
188	83	179	209	186	215	211	158	159	75	20	169
189	97	165	34	10	168	134	11	31	62	22	148
190	168	191	193	158	227	178	143	182	166	35	190
205	174	155	252	236	231	149	178	228	43	91	234
190	216	116	149	236	187	65	150	79	35	218	241
190	234	147	169	227	210	127	103	35	101	255	224
190	214	173	66	103	143	96	50	2	109	249	215
187	196	235	75	1	81	47	9	6	217	255	211
183	202	237	145	0	0	12	108	200	136	243	236
196	206	123	207	177	121	123	200	175	13	96	218

Images are Numbers




167	153	174	168	160	152	129	151	172	161	155	156
165	182	163	74	76	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	187
206	169	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	37	71	201
172	165	207	233	233	214	220	239	228	98	74	206
188	83	179	209	186	215	211	158	159	75	20	169
189	97	165	34	10	168	134	11	31	62	22	148
190	168	191	193	158	227	178	143	182	166	35	190
205	174	155	252	236	231	149	178	228	43	91	234
190	216	116	149	236	187	65	156	79	33	218	241
190	234	147	168	227	210	127	103	35	101	255	224
190	214	173	66	103	143	96	50	2	109	249	215
187	196	235	75	1	81	47	9	6	217	255	211
183	202	237	145	0	0	12	108	200	136	243	236
196	206	123	207	177	121	123	200	175	13	96	218

What the computer sees

157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	76	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	169	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	251
172	165	207	233	233	214	220	239	228	98	74	206
188	83	179	209	186	215	211	158	159	158	139	20
189	97	165	34	10	168	134	11	31	62	22	148
190	168	191	193	158	227	178	143	182	166	35	190
205	174	155	252	236	231	149	178	228	43	91	234
190	216	116	149	236	187	65	156	79	33	218	241
190	224	147	168	227	210	127	103	35	101	255	224
190	214	173	66	103	143	96	50	2	109	249	215
187	196	235	75	1	81	47	9	6	217	255	211
183	202	237	145	0	0	12	108	200	136	243	236
196	206	123	207	177	121	123	200	175	13	96	218

An image is just a matrix of numbers [0,255]!
i.e., 1080x1080x3 for an RGB image

Tasks in Computer Vision



- **Regression:** output variable takes continuous value
- **Classification:** output variable takes class label. Can produce probability of belonging to a particular class

High Level Feature Detection

Let's identify key features in each image category



Nose,
Eyes,
Mouth



Wheels,
License Plate,
Headlights



Door,
Windows,
Steps

Manual Feature Extraction

Domain knowledge

Define features

Detect features
to classify

Problems?

Manual Feature Extraction

Domain knowledge

Define features

Detect features
to classify

Viewpoint variation



Scale variation



Deformation



Occlusion



Illumination conditions



Background clutter



Intra-class variation



Manual Feature Extraction

Domain knowledge

Define features

Detect features
to classify

Viewpoint variation



Scale variation




Deformation



Occlusion



Illumination conditions



Background clutter




Intra-class variation



Learning Feature Representations


Can we learn a **hierarchy of features** directly from the data instead of hand engineering?

Low level features



Edges, dark spots

Mid level features



Eyes, ears, nose

High level features




Facial structure

Learning Visual Features




Fully Connected Neural Network



Fully Connected Neural Network

Input:

- 2D image
- Vector of pixel values




Fully Connected:

- Connect neuron in hidden layer to all neurons in input layer
- No spatial information!
- And many, many parameters!

Fully Connected Neural Network

Input:

- 2D image
- Vector of pixel values




Fully Connected:

- Connect neuron in hidden layer to all neurons in input layer
- No spatial information!
- And many, many parameters!

How can we use **spatial structure** in the input to inform the architecture of the network?



Using Spatial Structure

Input: 2D image.
Array of pixel values




Idea: connect patches of input
to neurons in hidden layer.
Neuron connected to region of
input. Only "sees" these values.

Using Spatial Structure



Connect patch in input layer to a single neuron in subsequent layer.
Use a sliding window to define connections.
How can we **weight** the patch to detect particular features?

Feature Extraction with Convolution



- Filter of size 4×4 : 16 different weights
- Apply this same filter to 4×4 patches in input
- Shift by 2 pixels for next patch

This “patchy” operation is **convolution**

- 1) Apply a set of weights – a filter – to extract **local features**
- 2) Use **multiple filters** to extract different features
- 3) **Spatially share** parameters of each filter

Feature Extraction and Convolution

A Case Study

X or X?




-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	-1	-1	-1	-1	-1	-1	1	-1
-1	1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	1	-1	-1	1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

Image is represented as matrix of pixel values... and computers are literal!
We want to be able to classify an X as an X even if it's shifted, shrunk, rotated, deformed.

Features of X




Filters to Detect X Features

filters


$$\begin{matrix} 1 & -1 & -1 \\ -1 & 1 & -1 \\ -1 & -1 & 1 \end{matrix}$$

$$\begin{matrix} 1 & -1 & 1 \\ -1 & 1 & -1 \\ 1 & -1 & 1 \end{matrix}$$

$$\begin{matrix} -1 & -1 & 1 \\ -1 & 1 & -1 \\ 1 & -1 & -1 \end{matrix}$$




The Convolution Operation



The Convolution Operation


Suppose we want to compute the convolution of a 5x5 image and a 3x3 filter:



We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs...


The Convolution Operation

We slide the 3×3 filter over the input image, element-wise multiply, and add the outputs:




The Convolution Operation

We slide the 3×3 filter over the input image, element-wise multiply, and add the outputs:




The Convolution Operation

We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs:




The Convolution Operation

We slide the 3×3 filter over the input image, element-wise multiply, and add the outputs:




The Convolution Operation

We slide the 3×3 filter over the input image, element-wise multiply, and add the outputs:




The Convolution Operation

We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs:




The Convolution Operation

We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs:



The Convolution Operation


We slide the 3×3 filter over the input image, element-wise multiply, and add the outputs:



The Convolution Operation

We slide the 3x3 filter over the input image, element-wise multiply, and add the outputs:

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0



1	0	1
0	1	0
1	0	1

filter



4	3	4
2	4	3
2	3	4

feature map

Producing Feature Maps



Original



Sharpen




Edge Detect




“Strong” Edge
Detect

Feature Extraction with Convolution




- 1) Apply a set of weights – a filter – to extract **local features**
- 2) Use **multiple filters** to extract different features
- 3) **Spatially share** parameters of each filter

Convolutional Neural Networks (CNNs)



CNNs for Classification



`tf.keras.layers.Conv2D`

`tf.keras.activations.*`

`tf.keras.layers.MaxPool2D`

1. Convolution: Apply filters to generate feature maps.

2. Non-linearity: Often ReLU.

3. Pooling: Downsampling operation on each feature map.



`torch.nn.Conv2d`


`torch.nn.ReLU, ...`

`torch.nn.MaxPool2d`

Train model with image data.

Learn weights of filters in convolutional layers.

Convolutional Layers: Local Connectivity




`tf.keras.layers.Conv2D`

`torch.nn.Conv2d`

For a neuron in hidden layer:

- Take inputs from patch
- Compute weighted sum
- Apply bias

Convolutional Layers: Local Connectivity



4x4 filter: matrix
of weights w_{ij}

$$\sum_{i=1}^4 \sum_{j=1}^4 w_{ij} x_{i+p,j+q} + b$$

for neuron (p,q) in hidden layer

`tf.keras.layers.Conv2D`


`torch.nn.Conv2d`

For a neuron in hidden layer:

- Take inputs from patch
- Compute weighted sum
- Apply bias

- 1) applying a window of weights
- 2) computing linear combinations
- 3) activating with non-linear function

CNNs: Spatial Arrangement of Output Volume



Layer Dimensions:

$h \times w \times d$

where h and w are spatial dimensions
d (depth) = number of filters

Stride:

Filter step size

Receptive Field:

Locations in input image that
a node is path connected to



```
tf.keras.layers.Conv2D( filters=d, kernel_size=(h,w), strides=s )
```




```
torch.nn.Conv2d( in_channels=3, out_channels=d, kernel_size=(h,w), stride=s )
```


Need to specify
input dimensionality!

Introducing Non-Linearity

- Apply after every convolution operation (i.e., after convolutional layers)
- ReLU: pixel-by-pixel operation that replaces all negative values by zero. **Non-linear operation!**




Rectified Linear Unit (ReLU)



 `tf.keras.layers.ReLU`

 `torch.nn.ReLU`

Pooling




max pool with 2x2 filters
and stride 2

`tf.keras.layers.MaxPool2D(
 pool_size=(2, 2),
 strides=2
)`




`torch.nn.MaxPool2d(
 kernel_size=(2, 2),
 stride=2
)`




- 1) Reduced dimensionality
- 2) Spatial invariance

How else can we downsample and preserve spatial invariance?

Representation Learning in Deep CNNs




CNNs for Classification: Feature Learning



1. Learn features in input image through **convolution**
2. Introduce **non-linearity** through activation function (real-world data is non-linear!)
3. Reduce dimensionality and preserve spatial invariance with **pooling**

CNNs for Classification: Class Probabilities



- CONV and POOL layers output high-level features of input
- Fully connected layer uses these features for classifying input image
- Express output as **probability** of image belonging to a particular class


$$\text{softmax}(y_i) = \frac{e^{y_i}}{\sum_j e^{y_j}}$$

Putting it all together: CNN in TensorFlow



```
import tensorflow as tf

def generate_model():
    model = tf.keras.Sequential([
        # first convolutional layer
        tf.keras.layers.Conv2D(32, filter_size=3, activation='relu'),
        tf.keras.layers.MaxPool2D(pool_size=2, strides=2),
        # second convolutional layer
        tf.keras.layers.Conv2D(64, filter_size=3, activation='relu'),
        tf.keras.layers.MaxPool2D(pool_size=2, strides=2),
        # fully connected classifier
        tf.keras.layers.Flatten(),
        tf.keras.layers.Dense(1024, activation='relu'),
        tf.keras.layers.Dense(10, activation='softmax') # 10 outputs
    ])
    return model
```



Putting it all together: CNN in PyTorch




```
import torch

def generate_model():
    model = nn.Sequential([
        # first and second convolutional layer
        torch.nn.Conv2d(in_channels=3, out_channels=32, filter_size=3),
        torch.nn.ReLU(),
        torch.nn.MaxPool2d(kernel_size=2, stride=2),


        torch.nn.Conv2d(in_channels=32, out_channels=64, filter_size=3),
        torch.nn.ReLU(),
        torch.nn.MaxPool2d(kernel_size=2, stride=2),

        # fully connected classifier
        torch.nn.Flatten(),
        torch.nn.Linear(64*6*6, 1024), # flattened dim after 2 conv layers
        torch.nn.ReLU(),
        torch.nn.Linear(1024, 10), # 10 outputs
    ])
    return model
```



An Architecture for Many Applications


An Architecture for Many Applications






Classification: Breast Cancer Screening

International evaluation of an AI system for breast cancer screening

b



c





CNN-based system outperformed expert radiologists at detecting breast cancer from mammograms

Breast cancer case missed by radiologist but detected by AI

Object Detection



Image X





Taxi

Class label y



Image X





Label (x, y, w, h)

Object Detection



Image X





Output:

taxi: (x_l, y_l, w_l, h_l)



Image X



Output:



taxi: (x_l, y_l, w_l, h_l)

person: (x_2, y_2, w_2, h_2)

person: (x_3, y_3, w_3, h_3)

...

Naïve Solution to Object Detection




Problem: Way too many inputs! This results in too many scales, positions, sizes!

Object Detection with R-CNNs

R-CNN algorithm: Find regions that we think have objects. Use CNN to classify.



1. Input image



2. Extract region proposals (~2k)

3. Compute CNN features

4. Classify regions

Problems: 1) Slow! Many regions; time intensive inference.
2) Brittle! Manually defined region proposals.

Faster R-CNN Learns Region Proposals


Feature extraction over
proposed regions

Image input directly into
convolutional feature extractor
Fast! Only input image once!

Classification of regions →
object detection

Region proposal network
to learn candidate regions

Learned, data-driven





Semantic Segmentation: Fully Convolutional Networks

FCN: Fully Convolutional Network.

Network designed with all convolutional layers,
with **downsampling** and **upsampling** operations




Input:
 $3 \times H \times W$



Predictions:
 $H \times W$



`tf.keras.layers.Conv2DTranspose`




`torch.nn.ConvTranspose2d`

Continuous Control: Navigation from Vision


Raw Perception
 I
(ex. camera)



Coarse Maps
 M
(ex. GPS)




Possible Control Commands



End-to-End Framework for Autonomous Navigation

Entire model is trained end-to-end **without any human labelling or annotations**






$$L = -\log(P(\theta|I, M))$$



Auto ON
Navigation and Localization



Deep Learning for Computer Vision: Impact



Deep Learning for Computer Vision: Summary


Foundations

- Why computer vision?
- Representing images
- Convolutions for feature extraction



CNNs

- CNN architecture
- Application to classification
- ImageNet



Applications

- Segmentation, image captioning, control
- Security, medicine, robotics





MIT 6.S191

Introduction to Deep Learning

Lab 2: Facial Detection Systems

Link to download labs:

introtodeeplearning.com#schedule
github.com/MITDeepLearning/introtodeeplearning

1. Open the lab in Google Colab
2. Start executing code blocks and filling in the #TODOs
3. Need help? Come to 32-123!



Introduction to Deep Learning

Announcements and Reminders

Lecture 4: Generative Modeling starting soon!

Lab 2 begins today!
Prizes for participants!
Enter for a chance to win \$\$\$!

Lots of swag will be available after
class tomorrow!

Interested in cutting-edge AI jobs?
introtodeeplearning.com/jobs.html